

CERTIFICATE OF ATTENDANCE



To acknowledge that

Hugo Oliveira

attended

Foundations of LLM Mastery: Fine-tuning on one GPU

A blue cursive signature of Claudia Blaas.

Dr. Claudia Blaas-Schenner
Head of Learning Center
AI Factory Austria AI:AT
Vienna, 24 March 2026

Course: <https://events.asc.ac.at/e/LLM-ft1-2026-03>
Level: Intermediate
Duration: 3 hours
Date: 24 March 2026

Certificate of attendance for Hugo Oliveira

Foundations of LLM Mastery: Fine-tuning on one GPU

Fine-tuning large language models (LLMs) often requires substantial resources, yet recent advancements allow for effective training on a single GPU, reducing both the memory footprint as well as the computational load.

This practical course introduced industry professionals – from start-ups to SMEs and large enterprises – to the key concepts of the transformer architecture as well as techniques for compressing LLMs and fine-tuning them efficiently without sacrificing model performance.

During the course participants developed a foundational understanding of techniques reducing the memory footprint of a large language model during fine-tuning and inference.

Topics covered:

- Transformer Anatomy
- Tokenization & Embeddings
- Hugging Face Ecosystem
- Hands-on exercises
- Quantization
- PEFT
- Unsloth

Course content level:

- Intermediate

Course format:

- Online lectures, demos, and hands-on labs of **3 hours** (breaks excluded)
- Equivalent to **0.125 ECTS** credits

Lectures taught by:

- Simeon Harrison (AI Factory Austria AI:AT)
- Martin Pfister (AI Factory Austria AI:AT)

Hands-on labs:

- Conducted on the Booster partition (A100 GPUs) of LEONARDO

Organized by:

- AI Factory Austria AI:AT

Course description:

- <https://events.asc.ac.at/e/LLM-ft1-2026-03>